

# Research Proposal: Computational resource management of multi channel controller

Hodai Goldman (Hodaig@cs.bgu.ac.il)

Department of Computer Science,  
Ben-Gurion University of the Negev, Beer Sheva, Israel

Supervisor: Dr. Gera Weiss

September 13, 2016

# 1 Background and Problem Statement

Today's computer power allows for consolidation of controllers towards systems where a single computer regulates many control loops, each with its varying needs of computation resources. This brings two research challenges that we intend to attack in the proposed thesis:

- How to schedule control tasks in order to achieve good performance in terms of control measures (overshoot, convergence speed, etc.)?
- What is a good interface for co-design of scheduling and control?

While it is possible to build control systems using standard operating systems, either real-time or desktop, with static or with dynamic scheduling schemes, there is an agreed opinion in the control community that these do not serve well for the purpose outlined above. For example, in [5], the authors say:

‘The delay and jitter introduced by the computer system can lead to significant performance degradation. To achieve good performance in systems with limited computer resources, the constraints of the implementation platform must be taken into account at design time.’

Similar views are expressed also in other papers [1, 11, 9].

Desktop type operating systems, like Windows and Linux, schedule for computational efficiency but do not allow for worst-case performance guarantees. Real-time operating systems, on the other hand, sacrifice some efficiency for timing predictability but the type of timing guarantees that such systems provide are not enough to be used for guaranteeing control performance. When using such operating systems for control, engineers usually apply controllers that work in a fixed periodic manner so that control behavior becomes deterministic and control performance can be guaranteed. This is not efficient because resources can be better utilized if controllers act at higher frequencies only when needed. In this work we will develop methods to combine the efficiency of desktop operating systems with the predictability of real-time operating systems in a way that is more suitable for control systems than periods and deadlines. We will show that applying control computations at non-periodic, carefully designed times, based on run-time data, allows for better utilization of the computational resources and for better control performance.




Figure 1: A typical close feedback control loop where the physical plant (*System*) is monitored via an array of sensors (*Sensing*) which produce noisy sample of the state variables. And after sensing, *State Estimator* aggregates the raw data from all the sensors and produce estimation of the current state. Then the *Control Law* entity uses actuators (not marked here for simplicity matters) in order to close the gap between the current estimated state and the reference state (*Input*).

The control loops that we will consider in the proposed thesis are of the architectural form shown in Figure 1. We assume a close feedback control loop where the physical plant, *System*, is monitored via an array of sensors (*Sensing*) which produce raw data that represents noisy sample of some functions of the state variables. We assume uncertainty observations from the sensors so, after sensing, an entity called *State Estimator* aggregates the raw data from all the sensors and makes an educated estimation of the current state. Then the *Control Law* entity uses actuators (devices such as engines or manipulators that change the state of the system) in order to close the gap between the current estimated state and the reference state (*Input*).

The current state of the art, as descibed e.g. in [?], is that control engineers first design control tasks as periodic computations, then they specify the required periodic frequency for the task, and then software engineers design a scheduler that ensures that the periodic frequency requirements are met. The last step is usually done using pre-computed knowledge of the expected (maximum) duration of the tasks. We claim that we can achieve better performance and better resource utilization for control systems by using richer and more flexible requirements for the tasks. Specifically, we will develop tools with which control engineers will be able to specify in a natural way features of their control loop that the scheduler will use for executing dynamic resource assignment that will, at the same time, guarantee required control performance and will be efficient in its use of computational resources.

## 2 Case studies

### 2.1 Vision based controllers for drones

To test our concepts, we will examine the implementation of an autonomously flying quad-rotor in the context of an agriculture case study. Specifically, we will implement a quad-rotor that flies in corridors and greenhouses by a vision based feedback. The challenge in this case study, from our perspective, is that image processing is a heavy computational task that requires careful scheduling in order to preserve the system predictability and stability.

The current state-of-the-art solution for involving heavy computational tasks such as vision in the control system is simply by adding computational power to the system (use faster processors), usually much more than needed, in order to eliminate the chance of losing predictability. Some, more conservative, control engineers prefer to isolate the heavy computation from the core control loop by allocating one of the processor's core for that task, or even run the vision processing on a different computer (APM, for example suggests to use image processing by adding a dedicated computer board [6]).

In Section 3 we propose an alternative framework for such systems, to be implemented in the proposed thesis, that allows to integrate the heavy computations in the control loop in an efficient way that, we believe, will allow for cutting the costs involved in adding a dedicated boards or in dedicating a computation core.

### 2.2 Nano-Satellite

Another case study we will examine is the control of a nano-satellite. Specifically, we will examine how our new approach can be applied in the context of scheduling the sensing tasks of IAI (Israel Aerospace Industries) nano satellites. Their satellites are controlled by a small, relatively slow, processor and the developers say that a central issue in programming the satellite is how to schedule sensing tasks. The current approach they are using is to periodically schedule all the sensing task every iteration, and they say that, because of sensor redundancy, this is the main computational consumption and is becoming to be too much for the capabilities of their processor.

We think that much of the sensing is unnecessary and that the controlling task can be achieved without "most up-to-date" information. Some times we can base on good estimation of the sensing value or even use the last value. As a secondary goal to our main goal of designing a good architecture for vision based drones (our main case study), we will also check the possibility to improve the nano-satellite issue with our framework.




Figure 2: Our proposal of a general scheduling framework. Each control loop (depicted in blue) informs the resource allocator (Scheduler) of its quality of estimation and the allocator allocates accordingly the computation resources among all the control loops. The underlying assumption here is that the noise in the sensed data is a function of the amount of computation resources. We assume that the more resources are invested in sensing the better (less noisy) sensed data is obtained.

### 3 Research Plan

Our research plan consists of: (1) Designing a methodology for effective allocation of computational resources in real-time control systems; and then (2) Demonstrating our methodology with a framework for developing such systems and with a case study. The details of these two steps are elaborated below.

#### 3.1 The proposed methodology

The general methodology that we will develop is illustrated in Figure 2. The methodology comes to support an efficient scheduling protocols in modern control systems consisting of a computer that runs many tasks that implement the control laws of independent control loops. The current state of art, as describrd above and we observed, e.g., in the code of the APM [6], is that the designers of each control loop specify a fixed rate for invocations of the corresponding control task. In our methodology, shown in Figure 2, there is a richer interaction between the *Scheduler* and the control loops. We suggest that each control loop (blue components) will tell the resource allocator (*scheduler*) of its level of certainty and that the allocator will base

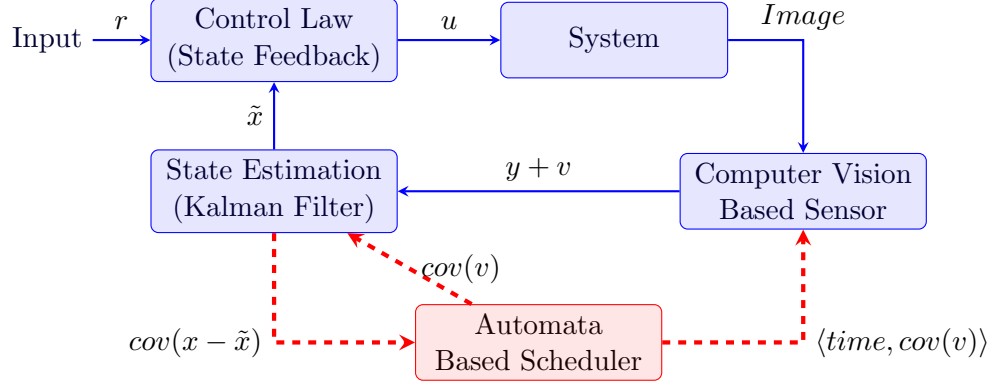


Figure 3: The controller framework we will implement, the *scheduler* will allocate CPU time ( $\langle time, cov(v) \rangle$ ) for the *Computer Vision* task base on the state estimation certainty using guarded Automata.

its allocation of CPU time based on this data. This can be useful on its own sake or together with a mathematical analysis that allows to use this feedback mechanism to certify some pre-defined specifications of state certainty or stability is maintained.

Our proposed methodology is general and may be applicable in a wide range of applications. However, in this initial phase of the research, we believe that it is better to focus on a specific sub-domain and in handling all technical issues in order to prove the concept. In this thesis we will develop and implement the, more specific, framework shown in Figure 3: a vision based controller for drone.

As shown in Figure 3, the suggested control system framework consists of an interaction between the tasks scheduler and the control loops, in this case the estimation task (*State Estimation*) accuracy strongly depends on the sensing task (*Computer Vision*) and they both collaborate with the *scheduler* in order to achieve their **control objectives** (g.e. stability). In this framework the scheduler is part of the control logic and therefore it can make scheduling decisions based on the current control state, the scheduling of *Computer Vision* task depends on the accuracy of state estimation. For example, if the vision is clear, the *Computer Vision* will produce good measurement of the *System* and therefore good accuracy will be achieved so the *Scheduler* can allocate less computation time to the heavy *Computer Vision* task and still remain stable while allocating more CPU to others control loops or to background tasks (like navigation).

We will use an existing open-source implementation of drone software as a basis for our experimentation. However, the above collaboration requires to re-adjust some parts of the control system, e.g. the *State Estimator* needs to work with variable error covariance of *Computer Vision* measurement, the *Computer Vision* needs to be able run within variable time limits, and the tasks pre-defined requirements need to be re-adjusted in order to define the relation between tasks such as *Computer Vision* and *State Estimator*. Below we will dive in each part of the above system (Figure 3) and explain how it will be adjusted (changed) and how we will achieve that adjustment in our demonstration.

### 3.1.1 Sensors (*Computer Vision Based Sensor*)

*Computer Vision Based Sensor* is the module that is responsible for taking a picture or a series of pictures and produce measurements of quantities such as speed and position relative to the environment. In our research we will use the *Computer Vision* in order to detect two dimensional movement in the camera surface (i.e., the drone speed). There are many such algorithms (called optical flow) differing in running times and in accuracy. Usually more invested time leads to more accuracy.

In our case we need to be able to control the *Computer Vision* running time and to have some good knowledge of the solution accuracy in order to achieve optimal state estimation (see Section 3.1.2). We assume that the *Computer Vision* error is distributed normally, and we will use the error variance as a measure of accuracy. We believe that this is a reasonable assumption because the estimated speed is usually computed as the average of many independent random variables, as follows. Optic flow algorithms usually go by identifying similar regions in consecutive pictures and then averaging the distances that each feature “traveled”. Assuming that the error in measurement of each feature is independent of the other errors, we get that the total error is the average of independent random variables. Then, by the law of large numbers, we get that the error should have a normal distribution. We will validate this assumption by experimentation with different parameters of different algorithms.

In order to control the running time, we will use anytime based algorithms [11] and will mainly concentrate on “contract based” vision algorithms proposed by Pant [9]. This type of algorithms run until we stop them, and when we stop them they will provide a solution with accuracy that is a monotonic function of the amount of time it was running. That way we can control the solution accuracy by controlling the running time.

In our implementation, in order to lower the complexity, we will pre-define few specific “operation modes” of the *Computer Vision* task, that differ by their running time and they are identified by a pair  $\langle RunTime, Covariance \rangle$  where *Covariance* is the error covariance of running time *RunTime*.

### 3.1.2 State Estimator

In cases where we have uncertain observations (inaccurate sensors), e.g. the computer vision measurement, we prefer to make some **educated estimation** of the current state based on the uncertain observations rather than just consider the last observation as the current state. This **educated estimation** of the current state is the main goal of the *State Estimator* block. This block receives the raw data from the sensors and produce an estimation of the current state.

In this thesis we will use well known optimal estimator called Kalman filter, or, more generally, extended Kalman filter (EKF) which is the nonlinear version of Kalman filter. The algorithm works in a two-step process. In the prediction step, the Kalman filter produces estimates of the current state variables, along with their uncertainties. Once the outcome of the next measurement is observed, these estimates are updated using a weighted average, with more weight being given to estimates with higher certainty [8].

In order to produce optimal estimation with Kalman filter, one of the parameter that we need to consider in our calculation is the variance of the sensor error. But in our new framework the sensor (vision based sensor) have variable error variance for each time slot, this means that we need to have also variable state estimators correspondingly. In order to adjust the sensor error variance, each time step the scheduler will inform the state estimator about the new error variance, and the state estimator will use corresponding parameters to make the next estimation.

APM, the control software that we will use, is already using kalman filter as the state estimator, so we only need to adjust the existing module to the variable sensor error variance, and to send the estimation error variance to the scheduler (see Section 3.1.4).

### 3.1.3 Control Tasks

The control task itself (*Control Law* in Figure 2) is responsible for reducing the difference between the current state ( $x$ ) where  $\tilde{x} = x + \text{“estimation error”}$  (in Figure 2) and the target state ( $u$  in Figure 2), by manipulating the drone motors speed.



The control task is usually a very low CPU consumer and been well studied [3, 5]. Hence, in the proposed thesis, we will not focus on the control task and will use the existing controller that APM are using. APM use a commonly used and well known technique from control theory called PID, A proportional-integral-derivative controller. In this technique the control output is based on the physical knowledge of the system dynamics, and have three variable parameters (P, I and D) that define the controller convergence behavior.

We also plan to try to use an “adjustable” LQR (Linear-quadratic regulator) control, as follows. LQR is an optimal controller that take into account the level of certainty of the state estimation. Is is usually a hard task to know the certainty of estimation, but in our case we need to calculate it anyway, and we mark it as  $cov(x - \tilde{x})$ . We will make LQR “adjustable” in the sense that each iteration the controller will consider the new (variable)  $cov(x - \tilde{x})$ . This is similar to the architecture we proposed for the estimator. We will check if adjustable LQR has significant advantages over PID in cases were we have variable certainty of estimation.

#### **3.1.4 Computation Resource Scheduler: Automata Based Scheduler**

Real-time systems are mostly composed of multiple real-time tasks, tasks with time constraint, for example task that must response within specified time constraints. The purpose of schedulers in such systems is to allocate the limited computational resources (CPU time) within all the task in the system. To do so, we need a well defined interface between the real-time tasks and the scheduler.

The most common way of describing the requirements of a real-time component is to specify a period, sometimes along with a deadline, which gives the frequency at which the component must execute. The designer of the component makes sure that the performance objectives are met as long as the component is executed consistent with its period. The scheduler guarantees that all components get enough resources. Specifying resource requirements using periods has advantages due to simplicity and analyzability, but has limited expressiveness, as elaborated in [1].

In this thesis one of the main goals, as we said above in Section 3, is to develop a methodology for allocation resources. We will focus on proposing way for engineers to describe the requirements of a real-time component. To simplify, we will assume that the resource is allocated in discrete slots of some fixed duration in the style of time-triggered architecture [1]. We will

develop automata based (hybrid automata) methodologies and schedulers based on the work of Alur [1] and Bukra [4], and will demonstrate with measurable data how they improve the performance of real flying drones (see Section 4).

We will use hybrid automata as specification framework for resource requirements. Hybrid automata is a variation of finite automata for real-valued continuously progressing words [10]. Automata can be more expressive for describing specific requirements, and they are composable, i.e., it is easy to compose all the tasks requirements into an integrated automaton. In our setting, the hybrid automata define continuous (fixed period) operation modes, represented by the automata state, and the condition of mode changing (mode transition) are denoted by the edges in the automata. The operation mode directly define a set of tasks that will be executed in the time slot (for example  $s_{0.7}$  in Figure 4). We can stay at a single mode for some iterations, in this case the same tasks will be scheduled in each iteration. We can also take **discrete mode transition** in order to change operation mode after few iterations and schedule different set of tasks, for example, in Figure 4 the edge  $(m_1, m_2)$  says that if the previous iteration mode was  $m_1$  and  $estCov < 0.7$  we can change the mode and schedule the set  $\{s_{0.2}\}$  in the next iteration.

Each infinite path on the composed hybrid automata (not necessarily with infinitely many mode transitions) represents a schedule that satisfy all the tasks requirements. In this architecture the scheduler only need to “walk through” the composed automata, this is, of course, a fast and easy computational task. In order to assure that we do not exceed the time slot duration, each task will have pre-defined maximum duration time like “deadline” in the traditional architecture. Now we can verify that every possible scheduling step can execute within a single time slot, in other words, every mode of the composed automata can be executed in a single time slot, by simply summing the “deadlines” of all the tasks in the mode tasks set. If we find a mode that goes beyond the maximum duration we can remove it from the automata so we never exceed the time slot duration.

Let us demonstrate the proposed automata based interface using the system depicted in Figure 3. Assume we have two operation modes of the vision based sensor: (1)  $s_{0.7}$  a very accurate operation mode that takes 70% of the time slot to execute, that is of course a significant amount of time, and (2)  $s_{0.2}$  a less accurate operation mode but takes only 20% of the slot. In each time slot we can execute one of them and get the sensing process done, but if we use only  $s_{0.7}$  every time slot we may not have enough time to execute all the tasks. On the other hand, if we use only  $s_{0.2}$  we will have




Figure 4: Example of guarded automata for our vision based sensor task, in this example the task has two operation modes,  $\mathbf{s}_{0.7}$  which need 70% of the time slot but is more accurate vision computation and  $\mathbf{s}_{0.2}$  which is less accurate but faster (need only 20% of the slot). The value of  $estCov$  is  $cov(x - \tilde{x})$  of the previous iteration. Every time slot exactly one of the modes will be executed.




Figure 5: Example of guarded automata for our state estimator, in this example the estimator has two operation modes,  $\mathbf{est}_1$  correspond to  $\mathbf{s}_{0.7}$  and  $\mathbf{est}_2$  correspond to  $\mathbf{s}_{0.2}$  (see Figure 4).

inferior estimates. Figure 4 shows an example of a guarded automata that guides the schedule. With this automata we can express rich specifications. In this example, execute  $s_{0.7}$  is always allowed but if we need faster sensing we can use  $s_{0.2}$  but only once in a row if the estimation error is not extremely bad (if  $cov(x - \tilde{x}) < 0.7$ ), or even twice in a row if the estimation error is good (if  $cov(x - \tilde{x}) < 0.3$ ).

The estimated estimation error ( $estCov$  in the figure) is, in this case, the estimation error variance ( $cov(x - \tilde{x})$  in Figure 3).

This value ( $estCov = cov(x - \tilde{x})$ ) is calculated by the state estimator task (Section 3.1.2) and is passed to the scheduler as discussed before.

Each of this operation modes ( $s_{0.7}$  and  $s_{0.2}$ ) have different accuracy, specified by  $cov(v)$  in Figure 3, and if we want to get optimal estimation the state estimator must be configure correspondingly, i.e., the sensing error variance should be adjusted to the correct value ( $cov(s_{0.7})$ ), an easy solution for specifying the different configurations is by the guarded automata shown in Figure 5, which defines two operation modes of the state estimator,  $est_1$  and  $est_2$ , that correspond to  $cov(s_{0.7})$  and  $cov(s_{0.2})$ . So of course if we sense in mode  $s_{0.7}$  we must estimate with mode  $est_1$  that has the correct configurations for  $s_{0.7}$ , and if we sense in mode  $s_{0.2}$  we must estimate with mode  $est_2$ .

## 4 Preliminary Work and Results

In this preliminary phase we searched for the appropriate environment for experimentation. After research we decided to use Raspberry Pi with navio [7] board based drone with the commonly used and open-source controller software APM [6]. We managed to build and fly the complete drone, we dived into the APM code, we understand the code structure, and we are familiar with the relevant parts that we want to change.

We now have clear picture of our research plan for the near future in order to get basic results with the new schedule frame fork. In summary, our research plan is describe in Table 4.

As part of the preliminary work we developed (with two undergraduate students, as part of their final project in software engineering) the **Simu-Copter** [2] framework that allows to program our drone (or any other APM based drones) directly by Matlab Simulink model, using a code generation logic. We will add to that framework interfaces and tools for using our new scheduler capabilities embedded in the simulink diagrams, and give the control engineers a complete tool for developing control software. We also

Task	Expected work time	completion (%)
Choose the appropriate development environment. (APM based drones)	one month	<b>100%</b>
Be familiar with the development environment, and make sure that we are able to adapt the system to our needs.	three month	<b>90%</b>
Plan all the control components in order to achieve our case study goals. (see Sections 3.1.2, 3.1.1 and 3.1.3)	three months	<b>100%</b>
Plan the new scheduler detail behavior, and its interaction with the control components. (see Section 3.1.4)	two months	<b>90%</b>
Detail planing of the scheduler interface implementation and adaptation in the code.	one month	<b>80%</b>
Implement the new scheduler and replace the original APM scheduler (manually transform each task periodical requirements to automata based requirements).	two weeks	0%
Develop and implement the vision based sensor component. We plan to base on the built-in Matlab vision abilities and use SimuCopter [2] to integrate it in the APM program.	1-2 months.	20%
Grub statistics data and fine-tuning the vision component.	one month	0%
Adjust the state estimator as explain in Section 3.1.2, and combine all together using the new scheduler.	two weeks	0%
Check with a wide range of environmental conditions and a wide range of requirements (automatas), and of course grub and summarize the results.	two months	0%
Write the thesis report.	two months	10%

Table 1: Research plan

experimented with implementing computer vision algorithms with this tool.

## References

- [1] Rajeev Alur and Gera Weiss. Rtcomposer: A framework for real-time components with scheduling interfaces. *8th ACM/IEEE Conference on Embedded Software*, 2008.
- [2] Sagi Avasker Amit Bedarshi and Gera Weiss. SimuCopter: open source tool for programming quad-rotors with simulink. <https://github.com/flyngPig>, 2016.
- [3] Stuart Bennett. A brief history of automatic control. *IEEE Control Systems Magazine*, 16:17–25, 1996.
- [4] Merav Bukra. GameComposer: A Framework for Dynamic Scheduling. Master’s thesis, Ben-Gurion University, Israel, October 2014.
- [5] Anton Cervin, Dan Henriksson, Bo Lincoln, Johan Eker, and Karl-Erik Arzen. How does control timing affect performance?, 2003.
- [6] Ardupilot Developers. ArduPilot Mega (APM): open source controller for quad-copters. [www.ardupilot.org](http://www.ardupilot.org).
- [7] EMLID. navio: Linux autopilot on Raspberry Pi. <https://emlid.com/>.
- [8] Rudolph Emil Kalman. A new approach to linear filtering and prediction problems. *Transactions of the ASME–Journal of Basic Engineering*, 82(Series D):35–45, 1960.
- [9] Yash Vardhan Pant, Kartik Mohta, Houssam Abbas, Truong X. Nghiem, Joseph Devietti, and Rahul Mangharam. Co-design of anytime computation and robust control (supplemental). Technical Report UPenn-ESE-15-324, Department of Electrical and Systems Engineering, University of Pennsylvania, Philadelphia, PA, May 2015.
- [10] A.J. van der Schaft and J.M. Schumacher. *An Introduction to Hybrid Dynamical Systems*. Lecture Notes in Control and Information Sciences. Springer, 2000.
- [11] Shlomo Zilberstein and Stuart J. Russell. Optimal composition of real-time systems. *Artificial Intelligence*, 82(1-2):181–213, 1996.